


Binary search trees (BST)


Tree

- A tree is a structure that represents a parent-child relation on a set of object.
- An element of a tree is called a **node** or **vertex**.
- The **root** of a tree is the unique node that does not have a parent (node A in the example below).
- A node is a **leaf** if it doesn't have children (C,F,G,H,I).
- A node which is not a leaf is called an **inner node**.




Tree

- A node w is a **descendant** of a node v if there are nodes v_1, v_2, \dots, v_t , with $v_1 = v$ and $v_t = w$, such that v_{i+1} is a child of v_i for all i . v is called an **ancestor** of v .
- The **subtree rooted at v** is a tree containing v and its descendants.




The subtree rooted at B




Depth

- The **depth** of a node v is the number of edges in the path from the root to v .
In the example below, the depth of B is 1.




Height

- The **height** of a node v is the maximum number of edges in a path from v to a descendant of v .
In the example below, the height of B is 2.
- The **height** of a tree is the height of its root = the maximum depth of a leaf.
- The height of an empty tree is -1 .




Degree

- The **degree** of a node v is the number of children of v .
For example, $\text{degree}(A) = 3$.




Binary tree

- A **binary tree** is a tree in which each node has at most two children.
- More precisely, a node can have a **left child** or not, and have a **right child** or not.




Binary




Not binary

Full binary tree

- A **full binary tree** is a binary tree in which each internal node has **exactly** two children.




Full




Not full

Perfect binary tree

- A **perfect binary tree** is a full binary tree in which all leaves have the same depth.




Perfect



Not perfect

Complete binary tree

- A **complete binary tree** is either a perfect binary tree, or a tree that is obtained from a perfect binary tree by deleting consecutive leaves of the tree from right to left.



Dynamic set ADT

A **dynamic set ADT** is a structure that stores a set of elements. Each element has a (unique) **key** and **satellite data**. The structure supports the following operations.

Search(S, k) Return the element whose key is k .

Insert(S, x) Add x to S .

Delete(S, x) Remove x from S (the operation receives a pointer to x).

Minimum(S) Return the element in S with smallest key.

Maximum(S) Return the element in S with largest key.

Successor(S, x) Return the element in S with smallest key that is larger than $x.key$.

Predecessor(S, x) Return the element in S with largest key that is smaller than $x.key$.


Simple implementation of dynamic set ADT

- The dynamic set ADT can be implemented using linked list.
- Is this implementation good?

Binary Search Tree

A **binary search tree** (BST) is an implementation of the dynamic set ADT. The elements of the set are stored in the nodes of a binary tree (exactly one element in each node) such that the following property is satisfied for every node x .

- For every node y in the left subtree of x , $y.\text{key} \leq x.\text{key}$.
- For every node z in the right subtree of x , $z.\text{key} > x.\text{key}$.



Binary Search Tree

Each node x in the tree has the following fields.

- key: The key of the element of x .
- Satellite data.
- left: pointer to the left child of x (can be NULL).
- right: pointer to the right child of x (can be NULL)
- p: pointer to the parent of x (NULL for the root).

T .root is a pointer to the root of tree T .

Tree walk


INORDER(x)

- (1) **if** $x \neq \text{NULL}$
- (2) INORDER($x.\text{left}$)
- (3) **print** $x.\text{key}$
- (4) INORDER($x.\text{right}$)

Time complexity: $\Theta(n)$, where n is the number of nodes.
We also define two additional procedures.

Preorder Print $x.\text{key}$ before the two recursive calls.

Postorder Print $x.\text{key}$ after the two recursive calls.



Inorder: 2 3 5 6 7 8


Preorder: 6 3 2 5 8 7

Postorder: 2 5 3 7 8 6

Recursive search

SEARCH(x, k)

- (1) **if** $x = \text{NULL}$ OR $k = x.\text{key}$
- (2) **return** x
- (3) **if** $k < x.\text{key}$
- (4) **return** **SEARCH**($x.\text{left}, k$)
- (5) **else**
- (6) **return** **SEARCH**($x.\text{right}, k$)




Time complexity: $\Theta(h)$ where h is the height of the tree.

Recursive search

$\text{SEARCH}(x, k)$

- (1) **if** $x = \text{NULL}$ OR $k = x.\text{key}$
- (2) **return** x
- (3) **if** $k < x.\text{key}$
- (4) **return** $\text{SEARCH}(x.\text{left}, k)$
- (5) **else**
- (6) **return** $\text{SEARCH}(x.\text{right}, k)$




Time complexity: $\Theta(h)$ where h is the height of the tree.

Recursive search

$\text{SEARCH}(x, k)$

- (1) **if** $x = \text{NULL}$ OR $k = x.\text{key}$
- (2) **return** x
- (3) **if** $k < x.\text{key}$
- (4) **return** $\text{SEARCH}(x.\text{left}, k)$
- (5) **else**
- (6) **return** $\text{SEARCH}(x.\text{right}, k)$



Time complexity: $\Theta(h)$ where h is the height of the tree.

Iterative search

SEARCH(x, k)

```
(1) while  $x \neq \text{NULL}$  AND  $k \neq x.\text{key}$ 
(2)   if  $k < x.\text{key}$ 
(3)      $x \leftarrow x.\text{left}$ 
(4)   else
(5)      $x \leftarrow x.\text{right}$ 
(6) return  $x$ 
```

Time complexity: $\Theta(h)$.

Minimum/Maximum


MINIMUM(x)

- (1) **while** $x.\text{left} \neq \text{NULL}$
- (2) $x \leftarrow x.\text{left}$
- (3) **return** x

MAXIMUM(x)

- (1) **while** $x.\text{right} \neq \text{NULL}$
- (2) $x \leftarrow x.\text{right}$
- (3) **return** x

Time complexity: $\Theta(h)$.



Successor


To find the successor of x :

- If x has a right child, the successor is the node with minimum key in the subtree of the right child.
- Otherwise, the successor is the lowest ancestor of x whose left child is also an ancestor of x (if no such ancestor exists, the successor is NULL).

Time complexity: $\Theta(h)$.

SUCCESSOR(x)

- (1) **if** $x.\text{right} \neq \text{NULL}$
- (2) **return** MINIMUM($x.\text{right}$)
- (3) $y \leftarrow x.p$
- (4) **while** $y \neq \text{NULL}$ AND $x = y.\text{right}$
- (5) $x \leftarrow y$
- (6) $y \leftarrow y.p$
- (7) **return** y



Successor


To find the successor of x :

- If x has a right child, the successor is the node with minimum key in the subtree of the right child.
- Otherwise, the successor is the lowest ancestor of x whose left child is also an ancestor of x (if no such ancestor exists, the successor is NULL).

Time complexity: $\Theta(h)$.

SUCCESSOR(x)

- (1) **if** $x.\text{right} \neq \text{NULL}$
- (2) **return** MINIMUM($x.\text{right}$)
- (3) $y \leftarrow x.p$
- (4) **while** $y \neq \text{NULL}$ AND $x = y.\text{right}$
- (5) $x \leftarrow y$
- (6) $y \leftarrow y.p$
- (7) **return** y



Insertion

- When inserting a new element x to the tree, it is inserted as a leaf.
- The location of the new leaf is chosen such that the binary search tree property is maintained.
- This is done by going down from the root using the same algorithm as in the search procedure.
- We need to keep a “trailing pointer”, namely a pointer that points to the previous node visited by the algorithm.
- Time complexity: $\Theta(h)$.

Insertion

INSERT(T, z)

$y \leftarrow \text{NULL}$

$x \leftarrow T.\text{root}$

while $x \neq \text{NULL}$

$y \leftarrow x$

if $z.\text{key} < x.\text{key}$

$x \leftarrow x.\text{left}$

else

$x \leftarrow x.\text{right}$

$z.p \leftarrow y$

if $y = \text{NULL}$


$T.\text{root} \leftarrow z$ // T was empty

else if $z.\text{key} < y.\text{key}$

$y.\text{left} \leftarrow z$

else

$y.\text{right} \leftarrow z$



Insertion

INSERT(T, z)

$y \leftarrow \text{NULL}$

$x \leftarrow T.\text{root}$

while $x \neq \text{NULL}$

$y \leftarrow x$

if $z.\text{key} < x.\text{key}$

$x \leftarrow x.\text{left}$

else

$x \leftarrow x.\text{right}$

$z.p \leftarrow y$

if $y = \text{NULL}$


$T.\text{root} \leftarrow z$ // T was empty

else if $z.\text{key} < y.\text{key}$

$y.\text{left} \leftarrow z$

else

$y.\text{right} \leftarrow z$



Insertion

INSERT(T, z)

$y \leftarrow \text{NULL}$

$x \leftarrow T.\text{root}$

while $x \neq \text{NULL}$

$y \leftarrow x$

if $z.\text{key} < x.\text{key}$

$x \leftarrow x.\text{left}$

else

$x \leftarrow x.\text{right}$

$z.p \leftarrow y$

if $y = \text{NULL}$


$T.\text{root} \leftarrow z$ // T was empty

else if $z.\text{key} < y.\text{key}$

$y.\text{left} \leftarrow z$

else

$y.\text{right} \leftarrow z$



Insertion

INSERT(T, z)

$y \leftarrow \text{NULL}$

$x \leftarrow T.\text{root}$

while $x \neq \text{NULL}$

$y \leftarrow x$

if $z.\text{key} < x.\text{key}$

$x \leftarrow x.\text{left}$

else

$x \leftarrow x.\text{right}$

$z.p \leftarrow y$

if $y = \text{NULL}$


$T.\text{root} \leftarrow z$ // T was empty

else if $z.\text{key} < y.\text{key}$

$y.\text{left} \leftarrow z$

else

$y.\text{right} \leftarrow z$



Insertion

INSERT(T, z)

$y \leftarrow \text{NULL}$

$x \leftarrow T.\text{root}$

while $x \neq \text{NULL}$

$y \leftarrow x$

if $z.\text{key} < x.\text{key}$

$x \leftarrow x.\text{left}$

else

$x \leftarrow x.\text{right}$

$z.p \leftarrow y$

if $y = \text{NULL}$


$T.\text{root} \leftarrow z$ // T was empty

else if $z.\text{key} < y.\text{key}$

$y.\text{left} \leftarrow z$

else

$y.\text{right} \leftarrow z$



Deletion

To delete a node z in a binary search tree, there are 3 cases to consider:

Case 1 z has no children.


Case 2 z has 1 child.

Case 3 z has 2 children.

Time complexity: $\Theta(h)$.


Case 1

If z is a leaf, delete it from the tree.



Case 2


If z has a single child x , delete z from the tree, and make x a child of the parent of z .



Case 3


If z has two children:

- $y \leftarrow \text{SUCCESSOR}(z)$
- Delete y from the tree (using Case 1 or 2).
- Make y take the place of z in the tree.



Summary

- All the dynamic set operation take $\Theta(h)$ time (worst case) on a tree of height h .
- In the worst case, $h = n - 1$.
- In the best case, $h = \lfloor \log n \rfloor$ (in an optimal tree, level i of the tree contains exactly 2^i nodes, except perhaps the last level).



Randomly built binary search trees

A **randomly built binary search tree** is a tree that is built by inserting the elements in random order into an initially empty tree (each of $n!$ permutations of the elements is equally likely).

Theorem

The expected height of a randomly built binary search tree is $\Theta(\log n)$.

Little is known about the average height of a binary search tree when both insertion and deletion are used to create it.